PM COMPASS

Product Requirements Document

AI-Powered Knowledge Assistant for Product Managers

Field	Details
Product Name	PM Compass
Version	v1.0 — MVP
Date	March 2026
Author	Nagaswaroop Yelleped
Status	COMPLETE — All sections locked

Legend: LOCKED = content agreed | IN PROGRESS = being built | PLACEHOLDER = not yet started

---

PART 1 — TRADITIONAL PRD SECTIONS

---

1. Executive Summary

STATUS: LOCKED

PM Compass is an AI-powered knowledge assistant built exclusively for Product Managers at early-stage startups. PMs lose 35–50 minutes daily searching for documents and gathering meeting context — time spent navigating Google Drive by filename, opening results one by one, and piecing together context that should surface automatically. Existing tools like Glean solve parts of this problem but carry enterprise price tags ($60K minimum contract) that are out of reach for lean startup teams.

PM Compass connects to Google Workspace and lets PMs search all their documents by content, not filename, and walk into every meeting with a pre-generated brief built from their own documents. It uses a RAG architecture — retrieval-augmented generation — to answer natural language queries with cited, grounded responses, sourced directly from the PM's Google Drive.

The MVP has been built and validated. Core feasibility gates passed: ingestion pipeline, hybrid search, Cohere reranking, and GPT-4o-mini generation are all live at https://pm-compass.vercel.app. Google Drive integration is complete — 14 real documents indexed, clickable source citations in the UI. A two-layer eval framework is in place — retrieval evals (Hit Rate = 83.6%, F1 = 0.470) and LLM-as-judge (Faithfulness 4.89/5) — both launch gates now passed.

The target market is 47,000 PMs at 50–100 employee startups on Google Workspace, representing a $14.2M SOM. At $25/month, 1,500 users = $37,500 MRR within 6 months of launch, acquired through community-led GTM in PM communities.

Current status: MVP complete. Both launch gates passed — search relevance 83.6% (gate: ≥ 80%) and hallucination rate ≤ 10% (Faithfulness 4.89/5). Ready for beta.

---

2. Context of Market & Opportunity

STATUS: LOCKED

2.1 Market Sizing

Level	Size	Logic	Value
TAM	1.5M PMs globally	All PMs worldwide × $300/year	$450M
SAM	750K PMs on Google Workspace	TAM × 50% Google Workspace adoption	$225M
SOM	~47K PMs at 50–100 employee startups	1.5M × 5% sub-100 employee firms × 50% Google Workspace	$14.2M

2.2 Why Now — Three Converging Forces

- LLM APIs now cheap enough to build PM-specific tools at sub-$25/user pricing
- Google Workspace APIs mature and production-ready — OAuth, Drive, and Calendar integrations no longer experimental
- AI fatigue creating demand for role-specific tools: PMs want something that understands their workflow, not a generic assistant

2.3 Competitive Landscape

- **Glean:** $7.2B valuation, $200M ARR, PM-specific agents launched 2025. Gap: $60K+ minimum contract, requires IT team and months of implementation
- **Notion AI:** works only if team is Notion-first — most startups use Google Drive as primary doc infrastructure
- **Google Gemini for Workspace:** built for everyone, not PM-specific — no calendar-to-document intelligence

---

3. Strategic Alignment

STATUS: LOCKED

Acquire 1,500 paying users within 6 months of launch through community-led GTM, validating product-market fit in the startup PM segment.

---

4. Customer Needs

STATUS: LOCKED

4.1 Target Segment & Problem Statement

PM Compass is built for Product Managers at early-stage startups with 50–100 employees that run on Google Workspace — using Google Drive to store PRDs, roadmaps, meeting notes, effort estimations, and presentations as their primary operational infrastructure.

These PMs struggle to stay on top of everything happening in the company — the product they're building, dependencies between teams, stakeholders, and the business. Existing tools like Glean attempt to solve parts of this problem but are not built with PMs specifically in mind and carry enterprise price tags ($50+/user, $60K minimum contract) that are out of reach for lean startup teams.

4.2 The Pain — Day in the Life

PMs operate in a high-context environment where staying informed is a core job requirement. Research indicates PMs lose 20–30 minutes daily searching for documents and 15–20 minutes per meeting gathering context. *(Note: sourced from internet research — to be validated with user interviews.)*

4.3 Current Workaround

When a PM needs to find a document today, they navigate to Google Drive and search by filename. If the filename is unknown — common for new team members or cross-functional documents — they must open each result one by one. There is no way to search by content, topic, or context.

4.4 Jobs to Be Done (MVP Scope)

- **JTBD 1:** Find documents instantly by describing what they contain, not what they are named
- **JTBD 2:** Synthesize and summarize information across documents so PMs are better informed before meetings and decisions

---

5. Proposed Product

STATUS: LOCKED

5.1 Product Overview

PM Compass is an AI-powered knowledge assistant that connects to Google Workspace to help Product Managers find information instantly and walk into meetings prepared. It uses natural language search across all document types in Google Drive, matches documents to upcoming calendar events, and synthesises information across sources into concise, cited summaries.

5.2 MVP Scope

Included in MVP ✅	Not Included — V2 ❌
Google Drive search (Docs, Sheets, Slides)	Slack integration
Natural language queries across document content	Jira integration
Google Calendar integration	Confluence integration
Meeting brief generation	Mixpanel integration
Confidence scoring & source citations	Gmail integration
Graceful handling of no-result scenarios	Team workspaces

5.3 Functional Requirements

#	Requirement	Priority	JTBD
FR1	The system must accept natural language queries and match them against document content, not just filenames	Must	JTBD 1
FR2	The system must return the top 3 most relevant results with a relevance score for each	Must	JTBD 1

FR3	The system must prompt the user for more information when no relevant results are found on the first attempt	Must	JTBD 1
FR4	The system must inform the user when a document is unavailable or deleted after a second failed search attempt	Must	JTBD 1
FR5	The system must connect with Google Calendar and surface upcoming meetings for the current day	Must	JTBD 2
FR6	The system must match meeting context (subject line, agenda) against indexed documents and surface the most relevant ones	Must	JTBD 2
FR7	The system must summarise information across multiple matched documents and present it as a pre-meeting brief	Must	JTBD 2
FR8	The system must notify the user when no documents match a meeting agenda and prompt them to add context before generating a brief	Must	JTBD 2
FR9	The system must notify the user when a calendar event has no agenda and request input before attempting to generate a brief	Must	JTBD 2

5.4 Non-Functional Requirements

Category	Requirement	Target
Performance	Search query response time from submission to results displayed	< 5 seconds
Performance	Meeting prep brief generation time	< 10 seconds
Scalability	Architecture must support future connectors (Jira, Slack, Confluence, Mixpanel) without requiring a rebuild of core search and retrieval logic	Design decision — pre-planned connector interface
Reliability	Minimum search relevance — proportion of queries returning correct documents in top 3 results	≥ 80%
Reliability	Maximum hallucination rate — proportion of responses containing information not present in source documents	≤ 10%
Availability	System uptime to support teams across time zones including weekend usage	99% uptime target
Performance at Scale	Search latency must remain within the 5-second target as document volume grows — requires vector index optimisation	< 5 seconds maintained at scale

---

6. Differentiation

STATUS: LOCKED

Unlike Glean, PM Compass offers the same core capabilities dedicated to PM use cases at less than $25 a month.

Glean serves enterprise IT teams and general knowledge workers. PM Compass is built exclusively for how PMs think, search, and prepare — every feature tuned to PM workflows at a price accessible without a procurement process.

---

7. Competitive Advantage

STATUS: LOCKED

Two compounding moats:

- **Ruthless focus on PM use cases** that Glean cannot justify deprioritising enterprise contracts for. Every product decision made through one lens: does this make a startup PM's day better?
- **Growing dataset of PM-specific search behaviour and meeting patterns** that makes the product smarter for PMs in ways a horizontal tool structurally cannot replicate.

Speed of iteration as structural advantage: Glean's enterprise feedback cycles are long. PM Compass ships enhancements based on direct daily feedback from a single persona.

---

8. Go-to-Market Strategy

STATUS: LOCKED

Pre-Launch Landing page with Stripe paywall marketed in PM communities (Lenny's Newsletter, Product Hunt, LinkedIn). Collect real payment intent — refunded if no launch. Social listening for messaging validation.

Launch Activate waitlist, direct outreach to highest-intent leads, announce in communities with strongest traction. Founders personally onboard first 20 users.

Post-Launch Word of mouth as primary growth lever. Monitor quality bars (relevance ≥80%, hallucination ≤10%, latency <5s). Iterate weekly based on PM feedback.

---

9. Risk Analysis — Traditional

STATUS: LOCKED

Risk Category	The Risk	Validation Method
Value Risk	PMs don't feel the pain of manual document search strongly enough to change their current behaviour. Status quo bias — Google Drive is good enough even if it is slow.	Smoke test: landing page with feature overview marketed through PM communities. Track signups, mailer interest, and willingness to pay via Stripe paywall before product is built.
Usability Risk	Two layers: (1) PMs cannot figure out how to use the product without instruction. (2) PMs do not trust the AI output enough to act on it, leading to abandonment.	Beta test with zero onboarding instructions. Task completion test: give users a search task and observe if they complete it naturally. Track first-session abandonment rate.

Feasibility Risk	The Google OAuth and Drive API integration may not work. Without document access the entire pipeline fails and there is no product.	Technical spike in Week 1: validate Google OAuth connection, Drive API document fetch, and readable content extraction before committing to full build.
Business Viability Risk	Pricing is the highest risk — if priced wrong, early-stage startups won't adopt. Secondary risk: company documents sent to a third-party LLM API is a data privacy concern requiring a legal agreement before launch.	Pricing: Stripe paywall on landing page — real payment is the only true validator. Privacy: review and sign LLM provider Data Processing Agreement; add data usage disclosure in onboarding.

PART 2 — AI-SPECIFIC PRD SECTIONS

A. Dataset & Chunking Strategy

STATUS: LOCKED

PM Compass uses a RAG (Retrieval-Augmented Generation) architecture. Rather than training or fine-tuning an LLM on company documents, PM Compass retrieves the most relevant document chunks at query time and passes them to the LLM as context. This means the LLM answers from the PM's actual documents — not generic training data — enabling fresh, cited, hallucination-resistant responses without expensive model retraining.

Two Pipelines

- **Ingestion Pipeline** (runs once, then on document changes): Documents → Chunks → Embeddings → Stored in Supabase pgvector
- **Query Pipeline** (runs on every search): User query → Embed query → Match against stored embeddings via hybrid search → Retrieve top chunks → LLM generates response with citations

Data Strategy Detail

Component	Decision	Rationale
Data Sources	Google Drive (Docs, Sheets, Slides) + Google Calendar metadata + query logs	Covers both JTBDs: Drive for document search, Calendar for meeting prep. Query logs used post-launch to improve retrieval quality over time.
Chunking Strategy	Paragraph-level chunks (~300–500 tokens) with ~15–20% overlap between adjacent chunks	PM documents (PRDs, roadmaps, meeting notes) contain paragraph-level context. Too small = loses meaning. Too large = introduces noise, reducing relevance score. Overlap prevents boundary context loss.
Overlap Mechanism	Last ~15–20% of each chunk repeated at the start of the next chunk	Ensures ideas that span chunk boundaries exist fully in at least one chunk. Prevents retrieval gaps at document section transitions.
Content Type Handling	Structured docs (PRDs, roadmaps) chunked by section. Meeting notes chunked by paragraph. To be validated post-spike.	Different content types have different natural boundary units. Structured docs have headers; meeting notes are more conversational.

Embedding	text-embedding-3-small (OpenAI)	Each chunk converted to a vector representing its semantic meaning. Same model used for both chunk ingestion and query embedding to ensure consistent vector space.
Vector Storage	Supabase pgvector — indexed for fast approximate nearest-neighbour search	Included in Supabase free tier. Compatible with our stack. Built-in index optimisation for latency at scale.
Retrieval Method	Hybrid search: semantic similarity (embedding vectors) + keyword matching (Postgres full-text search)	PMs search two ways: vague natural language ("what did we decide about pricing") AND exact terms ("Q3 OKRs"). Semantic catches meaning; keyword catches exact terms. Hybrid is the production standard.
Ingestion Trigger	Runs once on initial connection. Re-runs when a document is added or modified in Google Drive.	Keeps the index fresh without requiring manual re-indexing. New documents are searchable immediately after update.

Open Questions (Section A)

- What is the optimal chunk size for PM documents? Starting at 300–500 tokens — to be validated through beta relevance scores.
- Embedding model confirmed as text-embedding-3-small pending DPA (see Section B).

B. Model Selection Strategy

STATUS: LOCKED

B.1 Decisions

Component	Selected	Cost	Rationale
LLM	GPT-4o mini (OpenAI)	$0.15 / 1M input tokens, $0.60 / 1M output tokens	Cheapest capable model for PM document Q&A. Start here — step up to GPT-4o only if hallucination rate exceeds 10% bar in evals.
Embedding Model	text-embedding-3-small (OpenAI)	$0.02 / 1M tokens	Lowest cost embedding model from OpenAI. High quality for English text. Same provider as LLM — one DPA, one API key, simpler stack.
Vector DB	Supabase pgvector	$0 (free tier)	Already in stack. No additional provider. pgvector handles both storage and approximate nearest-neighbour search.
Retrieval	Hybrid: pgvector (semantic) + Postgres full-text search (keyword)	No additional cost	Supabase supports both vector similarity and full-text search natively — hybrid search without a separate BM25 service.
Reranking	Cohere rerank-v3.5	Free trial; ~$1 / 1K searches	Re-scores top-N candidates by true semantic relevance after initial retrieval. Improves

		(production)	precision — reduces noisy chunks passed to LLM. Build/buy decision: Cohere's rerank API takes one line vs 2–3 weeks to build in-house.

B.2 Cost Estimate at Development Scale

With $5 of OpenAI credits during development and testing:

- Ingesting 14 Google Drive documents (~5 chunks each, ~200 tokens/chunk): ~67 chunks = ~13,400 embedding tokens = effectively free
- 100 test searches (query embedding + ~2,000 tokens of chunks per query + ~500 token response): ~250,000 tokens = ~$0.04
- Cohere reranking: free trial key covers development usage (10 calls/min limit)
- $5 covers approximately 10,000 test searches at development scale — more than enough to build, test, and run evals

B.3 Privacy & DPA

OpenAI provides a Data Processing Agreement (DPA) covering zero data retention for API calls. Required steps before handling real PM documents:

- Enable zero data retention in OpenAI API settings
- Review and sign DPA via OpenAI's enterprise privacy portal
- Add data usage disclosure to onboarding flow (covered in Section F — AI Safety)

B.4 Upgrade Trigger

Stay on GPT-4o mini unless: hallucination rate exceeds 10% in beta evals, or response quality is consistently rated poor by beta users. If either condition is met, step up to GPT-4o ($2.50 / 1M input tokens) and re-run evals.

C. AI Feasibility Analysis

STATUS: LOCKED

Summary

The Week 1 technical spike validated all three feasibility gates. The core RAG architecture is proven and deployed to production at https://pm-compass.vercel.app.

Feasibility Gates — Results

Gate	Question	Result
Embedding	Can we embed PM documents and store them in pgvector?	PASSED
Retrieval	Can hybrid search return relevant chunks from natural language queries?	PASSED
Generation	Can GPT-4o-mini generate grounded, cited answers from retrieved chunks?	PASSED

What Was Validated

Ingestion Pipeline

- 14 Google Drive documents fetched via Drive export API and processed end-to-end
- 67 chunks stored in Supabase pgvector with metadata (source file, source name, chunk index, content)

- Chunk size: 300–500 tokens with 15–20% overlap — confirmed sufficient for paragraph-level PM context
- Embedding model: text-embedding-3-small — consistent, low-cost, same provider as LLM
- Idempotent ingestion: old chunks cleared before re-ingestion to prevent duplicates

Query Pipeline

- Natural language queries embedded at query time using same model
- Hybrid search (pgvector semantic + Postgres full-text keyword) fetching top 10 candidates
- Relevance threshold of 0.30 applied to combined score — filters irrelevant chunks before reranking
- Cohere rerank-v3.5 re-scores remaining candidates by true semantic relevance, returns top 3
- GPT-4o-mini generating cited, grounded responses with streaming output
- End-to-end latency: embedding ~300ms, search ~200ms, reranking ~300ms, LLM ~1,500ms (streaming masks perceived wait)

Negative Case Validation

- TC-010: Query for "Salesforce integration" correctly returns "I could not find a confirmed answer" — system does not hallucinate on absent topics
- System prompt hardened to distinguish confirmed decisions from speculative content

Known Feasibility Risks

Risk	Status	Mitigation
Google Drive OAuth	✅ RESOLVED	OAuth flow complete; refresh token stored; 14 Drive docs live in production
Real document formats (Docs, Sheets, Slides)	✅ RESOLVED	Drive export API converts all formats to plain text — no format-specific parsers required
Token limits on large documents	Mitigated	Chunking strategy handles this; overlap prevents boundary loss
Retrieval precision	In progress	Reranking shipped (Cohere rerank-v3.5); precision 0.343 — manual eval and tuning in progress

Feasibility Verdict

GREEN — MVP complete and both launch gates passed. All core RAG components validated in production including Google Drive integration and reranking layer.

D. Total Cost of Ownership (TCO)

STATUS: LOCKED

Development Phase (Current)

Component	Cost	Notes
OpenAI API (embeddings + LLM)	~$5 for all dev/testing	Covers ~10,000 test searches at current usage
Supabase	$0 (free tier)	500MB storage, sufficient for test data

Vercel	$0 (hobby tier)	Auto-deploys from GitHub, sufficient for prototype
GitHub	$0	Public repo
Total dev cost	**~$5**	

Per-Query Cost Estimate (Production)

Step	Model	Tokens (est.)	Cost
Query embedding	text-embedding-3-small	~50 tokens	$0.000001
LLM generation	GPT-4o-mini	~1,500 input + ~300 output	~$0.0004
Total per query			**~$0.0004**

Monthly Cost at Scale

Usage Level	Queries/Month	Monthly API Cost
Early beta (50 users, 10 searches/day)	15,000	~$6
Growth (500 users, 10 searches/day)	150,000	~$60
Scale (2,000 users, 10 searches/day)	600,000	~$240

Infrastructure Cost at Scale

Component	Free Tier	Paid Tier	Trigger
Supabase	500MB, 2 projects	$25/month (Pro)	>500MB storage or >50K rows
Vercel	Hobby (1 project)	$20/month (Pro)	Team collaboration or custom domains
Total infra at growth		**~$45/month**	

TCO Summary

Phase	Monthly Cost	Notes
Development	~$5 one-time	OpenAI credits only
Beta (50 users)	~$6 API + $0 infra	Free tiers sufficient
Growth (500 users)	~$60 API + $45 infra	Upgrade Supabase + Vercel
Scale (2,000 users)	~$240 API + $45 infra	~$285/month total

Revenue context: At $25/month per user, 500 users = $12,500 MRR. Infrastructure cost at that scale = $105/month (0.8% of revenue). Unit economics are strong — cost scales sub-linearly with users.

Cost Optimisation Levers (v2)

- Caching frequent queries — top 20% of queries likely repeat; cache responses to reduce LLM calls

- Batch ingestion during off-peak hours — no urgency cost for document processing
- Upgrade to GPT-4o only for meeting brief generation (higher value, justify higher cost)

E. Evals & Minimum Bar for Production

STATUS: LOCKED

Overview

PM Compass uses a two-layer evaluation framework: retrieval evals (offline, deterministic) and LLM-as-judge evals (generative, scalable). Both layers are required — neither alone is sufficient to assess system health.

```
Question → [RETRIEVAL] → Chunks → [GENERATION] → Answer
               ↑                        ↑
         Precision/Recall          LLM-as-Judge
         (retrieval quality)       (answer quality)
```

Layer 1 — Retrieval Evaluation (Offline)

Method: Golden dataset of 67 QA pairs auto-generated by GPT-4o-mini from source chunks (regenerated after Drive corpus migration). Each pair contains a question, expected chunk, and source file.

Metrics:

Metric	Definition
Hit Rate	Did the expected chunk appear in retrieved results?
Precision	Of all chunks retrieved, what fraction were relevant?
Recall	Of all relevant chunks, what fraction were retrieved?
F1	Harmonic mean of precision and recall

Threshold Tuning Results (v1 — 52-chunk test corpus, no reranking):

Threshold	Hit Rate	Precision	Recall	F1	Decision
0.50	28.8%	0.240	0.288	0.256	Too strict — relevant chunks filtered out
0.40	46.2%	0.346	0.462	0.381	Recall improved, precision barely moved
0.30	78.8%	0.365	0.788	0.481	Best F1 — selected as threshold

Current baseline (v2 — 67-chunk Drive corpus, with Cohere reranking):

Metric	Score
Hit Rate	83.6%
Avg Precision	0.343
Avg Recall	0.836

Avg F1	0.470

Why 0.30 was chosen: Relevant chunks cluster in the 0.35–0.39 score range due to similar-topic documents. A threshold of 0.50 filtered out correct chunks. At 0.30, recall jumped significantly while precision held. Precision is the north star metric for PM Compass — false positives (wrong sources shown to a PM) erode trust. However, at 0.40, precision did not meaningfully improve vs 0.30, making 0.30 the best F1 tradeoff for v1.

Reranking impact: Cohere rerank-v3.5 shipped in v2. Hit rate improved from 78.8% → 83.6% (launch gate passed). Precision at 0.343 — this is partially a measurement artifact (3 chunks returned, 1 expected, so max precision per hit = 0.33) and partially genuine noise. Manual eval and tuning in progress to diagnose remaining precision gaps.

Layer 2 — LLM-as-Judge Evaluation (Generative)

Method: For each question, run the full production pipeline (embed → retrieve → generate). Pass the question, retrieved chunks, and generated answer to GPT-4o-mini as a judge. Judge scores on two dimensions using a 1–5 rubric.

Metrics:

Metric	Definition	Score < 4 indicates
Faithfulness (1–5)	Are all claims in the answer directly supported by the retrieved chunks?	Hallucination
Answer Relevance (1–5)	Does the answer actually address the question asked?	Off-topic response

Results (20-question sample, run on 52-chunk corpus — initial baseline):

Metric	Score	Interpretation
Avg Faithfulness	4.89 / 5	Strong — system prompt effectively grounding responses
Avg Answer Relevance	5.00 / 5	High — note: inflated by self-evaluation bias (see limitations)
Questions evaluated	18 / 20	2 skipped due to no chunks passing threshold

Results (19-question sample, run on 67-chunk Drive corpus — current baseline):

Metric	Score	Interpretation
Avg Faithfulness	4.86 / 5	Strong — consistent with initial baseline
Avg Answer Relevance	5.0 / 5	Inflated — auto judge cannot detect attribution hallucinations (see Layer 3)
Questions evaluated	14 / 19	2 graceful failures, 3 timeouts excluded

Known limitations:

1. **Self-evaluation bias** — GPT-4o-mini judges answers it effectively wrote. In production, upgrade judge to GPT-4o for independent evaluation.

2. **Cannot detect attribution hallucinations** — judge scores answers as relevant even when the named person is wrong. Manual eval (Layer 3) is required for this.
3. **Lenient on "I don't know" answers** — judge scores graceful failures as 5/5 relevant, which masks retrieval failures. Retrieval eval catches this gap.

How the Two Layers Complement Each Other

Retrieval score	LLM judge score	Diagnosis	Action
High	High	System healthy end-to-end	Monitor, no action
Low	High	Retrieval weak, LLM recovering — fragile	Fix retrieval first (reranking)
High	Low	Good retrieval, LLM hallucinating	Tighten system prompt or upgrade model
Low	Low	Broken end-to-end	Fix retrieval first

Launch Gates

Metric	Minimum Bar	Current Status
Search relevance (hit rate)	≥ 80%	✅ 83.6% — gate passed (reranking shipped)
Hallucination rate	≤ 10%	✅ Faithfulness 4.89/5 — gate passed
Search latency	< 5 seconds	✅ ~2 seconds — gate passed
Meeting brief generation	< 10 seconds	Not yet built — Phase 3 scope

Layer 3 — Manual Eval (Human-in-the-loop)

Method: 19 test cases manually scored across 7 categories (Exact Lookup, Factual Extraction, Tabular Extraction, Paraphrased, Cross-Document, Negative, Vague). Each question run against the live pipeline. Answer and retrieved chunks evaluated by a human judge.

Results (Drive corpus, 19 test cases):

Metric	Score
Precision (word overlap)	67%
Recall	3.9 / 5
Faithfulness	4.1 / 5
Answer Relevance	3.1 / 5
Hallucination Rate	28% (5/19 questions)
Pass Rate	63% (12/19)
Partial Pass	16% (3/19)

Fail	21% (4/19)

Failure analysis: All 4 failures involve attribution hallucination — the correct document is retrieved but the model names the wrong person. Root cause: fixed-size chunking cuts across ownership context. Multiple people mentioned in the same chunk confuse the model on attribution. Graceful failures (system returned "I don't know") all scored Pass.

Manual vs Auto Eval Comparison

Both evaluations run on the same 19 questions against the live Drive corpus:

Metric	Auto (LLM Judge)	Manual (Human)	Gap
Faithfulness	4.86 / 5	4.50 / 5	+0.36 auto
Answer Relevance	5.0 / 5	3.71 / 5	+1.29 auto

Why the gap exists: LLM-as-judge inflates Answer Relevance by ~1.3 points because it detects topical relevance, not factual accuracy. It scores an answer as "relevant" even when the named person is wrong. Manual eval correctly penalised attribution hallucinations — in 4 cases where the judge gave AR=5, the human judge gave AR=1–2. For a trust-critical PM tool, manual eval is the authoritative signal.

Eval Roadmap

Phase	Addition	Status	Purpose
v1	Retrieval eval + LLM-as-judge (sample)	✅ Shipped	Establish baseline, identify layer failures
v2	Reranking layer (Cohere rerank-v3.5)	✅ Shipped	Improve hit rate — achieved 83.6%
v2	Manual eval on Drive corpus (19 test cases)	✅ Complete	Diagnosed attribution hallucination as primary failure mode
v2	Manual vs auto eval comparison	✅ Complete	Confirmed auto eval inflates AR by ~1.3 points on hallucination cases
v3	Semantic chunking	Planned	Fix attribution errors — keep ownership context with decisions
v3	Upgrade judge to GPT-4o	Planned	Remove self-evaluation bias
Post-launch	Langfuse observability	Planned	Log every query, chunks, answer, scores
Post-launch	Online eval (log query + answer triplets)	Planned	Monitor quality drift in production
Post-launch	Usage-based scoring	Planned	Track which answers users act on vs discard

---

F. AI Safety Strategy

STATUS: LOCKED

F.1 Functional Risks

Risk	Minimum Bar	Pre-Launch Validation	Mitigation
Latency — response too slow	< 5s search, < 10s briefs	Measure response time end-to-end during beta	Optimise API calls; batch processing; caching for repeated queries
Relevance — wrong documents returned	≥ 80% accuracy (8/10)	Manual scoring: out of 10 searches how many return correct documents in top 3	Hybrid search (keyword + semantic); synonym mapping for PM terms; feedback button
Hallucination — AI generates information not in source	≤ 10% fabrication (9/10 accurate)	Manual review: count outputs with information absent from source documents	Mandatory source citations with direct links; confidence scoring; display exact quoted text

F.2 Data & Privacy Risks

Risk	Mitigation	Pre-Launch Validation
Company documents sent to third-party LLM API — risk of data retention or use for model training	Sign Data Processing Agreement (DPA) with LLM provider confirming zero data retention; add data usage disclosure in onboarding	Review and sign DPA before launch; confirm zero retention policy in writing
Google OAuth misconfiguration — users accessing documents they should not see	Strict OAuth scope configuration; users can only access their own Drive	OAuth scope testing: verify each user sees only their own documents
Personal information in documents exposed unintentionally	Data usage disclosure in onboarding; users consent before indexing begins	Review onboarding consent flow before launch

F.3 User Trust Risks

Risk	Mitigation	Pre-Launch Validation
Under-trust — bad early experience causes permanent abandonment	First three searches must return useful results; onboarding includes a guided test query so first experience is successful	Track first-session abandonment rate in beta; if PM searches once and does not return, flag as under-trust signal
Over-trust — PM acts on AI output without verifying and makes a bad decision	Always show source citations and confidence scores; include 'verify before acting' UX copy on every result	Deliberate error test in beta: plant a known mistake in one result and observe whether beta users catch it or act on it blindly

---

G. Live Prototype

STATUS: COMPLETE

Live prototype available at: https://pm-compass.vercel.app

- Natural language document search across real Google Drive documents
- Cohere reranking — top 3 results re-scored by semantic relevance before LLM
- Source citations on all results with clickable "Open in Drive ↗" links
- Side panel — click any source to view matching passage and open the original document
- Streaming responses (word-by-word output, reducing perceived latency)
- Latency breakdown per pipeline stage (embedding ms, search ms, LLM ms)
- Graceful no-result handling with user-facing messaging

> *Live data: 14 Google Drive documents (NexusCRM synthetic PM artifacts), 67 chunks stored in Supabase pgvector. Google OAuth integration complete.*

H. Defence Against Big Tech

STATUS: LOCKED

PM Compass's defence is depth over breadth. Google could ship native Workspace AI tomorrow and bundle it free — but it would be built for everyone, which means optimised for no one. PM Compass is built exclusively for how PMs think, search, and prepare. Every feature, every prompt, every relevance signal tuned to PM workflows. Over time, behavioural data from PM-specific usage compounds into a personalisation layer that a horizontal platform structurally cannot replicate. Big tech buys breadth. We own depth.

METRICS FRAMEWORK

STATUS: LOCKED

Overview

PM Compass success is measured across four layers. Each layer answers a different question.

We defined a metrics hierarchy — one north star (Return Rate D7) with 5 supporting metrics across product, business, system, and model layers, each tied to a specific product decision.

Priority Metrics — Shortlist for Decision-Making

Layer	Metric	Type	Why it matters
Product	Return rate (D7) — % of users returning within 7 days	Lagging ⭐ North Star	Proves PM Compass delivers enough value that PMs come back
Product	Brief generation rate — % of sessions that include a meeting brief	Leading (JTBD 2: walk into meetings prepared)	Proves PMs are using the meeting prep feature, not just search
Product	Zero-result rate — % of queries returning no answer	Leading (JTBD 1: find documents by content)	High rate = search is failing; low rate = PMs are finding what they need

Business	MRR — total monthly recurring revenue	Lagging	Proves willingness to pay; primary business health signal
System	Search latency (p95) — 95th percentile response time	Leading	Affects every interaction; above 5s = abandonment
Model	Hallucination rate — % of answers with unsupported claims	Lagging	Non-negotiable trust signal; wrong information permanently erodes PM trust

---

Full Metrics Detail

Layer 1 — Product Metrics

Are users getting value? Are they coming back?

Metric	Definition	Target (6 months)	Why it matters
Return rate (D7)	% of users who return within 7 days of first use	≥ 40%	North star — habit formation signal
Brief generation rate	% of sessions that include a meeting brief generation	≥ 20%	JTBD 2 adoption
Zero-result rate	% of queries returning "I could not find an answer"	≤ 15%	High = knowledge base gap or retrieval failure
Daily Active Users (DAU)	Unique users running at least 1 search per day	300	Core engagement signal
Search queries per session	Avg number of searches per visit	≥ 3	Low number = user not finding value

---

Layer 2 — Business Metrics

Is this a viable business?

Metric	Definition	Target (6 months)	Why it matters
Monthly Recurring Revenue (MRR)	Total recurring subscription revenue	$37,500 (1,500 users × $25)	Primary business health signal
Conversion rate	% of free trial users converting to paid	≥ 25%	Validates willingness to pay
Monthly churn rate	% of paying users cancelling per month	≤ 3%	High churn = product isn't sticky
Customer Acquisition Cost (CAC)	Total sales/marketing cost ÷ new customers	< $50	Community-led GTM keeps this low
ARR	MRR × 12	$450,000	Investor-facing metric

Layer 3 — System Metrics

Is the system reliable and fast?

Metric	Definition	Target	Why it matters
Search latency (p95)	95th percentile end-to-end response time	< 5 seconds	NFR gate — above this, users abandon
API error rate	% of queries returning an error	< 1%	Reliability baseline
Uptime	% of time system is available	≥ 99.5%	SLA for paying customers
Ingestion latency	Time from document added to searchable	< 60 seconds	Freshness expectation

---

Layer 4 — Model Metrics

Is the AI component working correctly?

Metric	Definition	Target	Measured by
Hallucination rate	% of answers containing unsupported claims	≤ 10%	LLM-as-judge (faithfulness < 4/5)
Search relevance (hit rate)	% of queries where expected chunk is in top 3	≥ 80%	Retrieval eval (golden dataset)
Faithfulness score	Avg LLM-as-judge faithfulness rating	≥ 4.5 / 5	LLM-as-judge
Answer relevance score	Avg LLM-as-judge relevance rating	≥ 4.5 / 5	LLM-as-judge
Precision (retrieval)	Relevant chunks / total chunks retrieved	≥ 0.50 (v2 target)	Retrieval eval
F1 score (retrieval)	Harmonic mean of precision and recall	≥ 0.60 (v2 target)	Retrieval eval

---

Metrics Hierarchy for Decision-Making

When metrics conflict, use this hierarchy:

1. **Hallucination rate** — non-negotiable. Wrong information shown to a PM erodes trust permanently. Block launch if above 10%.
2. **Search relevance** — primary utility metric. If users can't find answers, there's no product.
3. **Return rate** — habit signal. One-time users don't build a business.
4. **MRR** — lagging indicator of everything above working together.

APPENDIX

A. Open Questions

- Which LLM provider offers the best cost-to-quality ratio with a DPA available? *(Resolved: OpenAI — Section B)*
- What chunk size and overlap works best for PM documents? *(Starting at 300–500 tokens — validate in beta)*
- What is the full eval framework and automation strategy? *(Resolved: Section E)*
- What is the pricing model for V1? *(Draft: $25/month — to be validated via Stripe paywall)*
- How do we handle documents the PM does not own — shared Drive files?
- How do we maintain search latency under 5 seconds as document volume scales?

B. Version History

Version	Date	Changes
v0.1	Feb 2026	Initial template created
v0.2	Feb 2026	Customer Needs locked; Traditional Risk Analysis locked; AI Safety Strategy added
v0.3	Feb 2026	AI Safety Strategy rewritten from first principles — all three risk sub-sections locked
v0.4	Feb 2026	Section 5 (Proposed Product) locked — product overview, MVP scope, 9 FRs, 7 NFRs
v0.5	Feb 2026	Sections 2, 3, 6, 7, 8, H locked — market sizing, strategic alignment, differentiation, GTM, defence against big tech
v0.6	Feb 2026	Section A (Dataset & Chunking Strategy) locked
v0.7	Mar 2026	Section B (Model Selection) locked — GPT-4o mini, text-embedding-3-small, OpenAI DPA confirmed
v0.8	Mar 2026	Sections C, D, E locked — feasibility spike results, TCO, full eval framework with retrieval evals + LLM-as-judge. Metrics Framework added.
v0.9	Mar 2026	Section 1 (Executive Summary) locked — PRD complete. All sections locked.
v1.0	Mar 2026	Updated with live results: Google Drive integration complete (14 docs, 67 chunks), Cohere reranking shipped, eval baseline updated (Hit Rate 83.6%, F1 0.470), both launch gates passed.
v1.1	Mar 2026	Manual eval complete (19 test cases, Drive corpus). Attribution hallucination identified as primary failure mode (28% rate). Manual vs auto eval comparison added — auto inflates AR by ~1.3 points. JTBD-2 (Summarize doc) shipped. Eval roadmap updated with v3 fixes.